


Message Passing Programming

- Most widely used approach
- Minimal requirements on the underlying hardware

Key characteristics

Partitioned
Address space

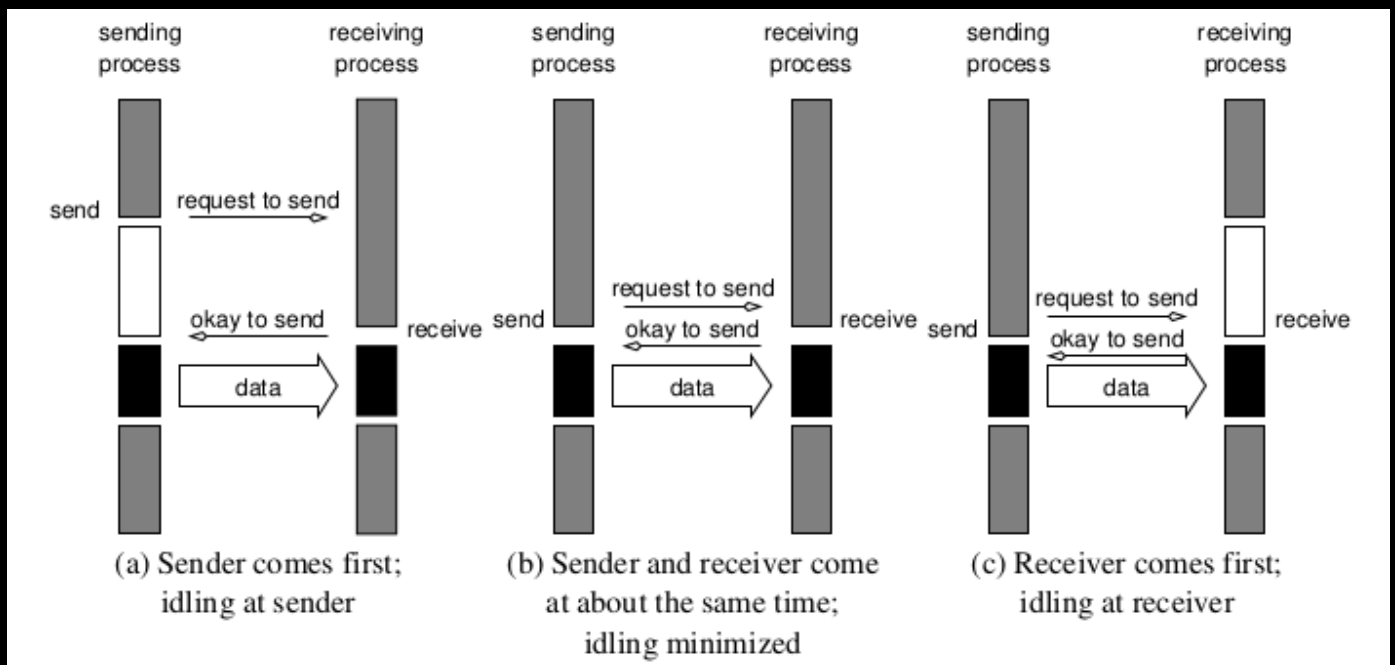
Explicit
Parallelization



- ⇒ Each data element must belong to one of the partitions
- ⇒ Needs explicit partitioning
- ⇒ Complexity in programming
- ⇒ But locality of access for high performance
- ⇒ All interactions need cooperation of two processes
- ⇒ Process that has the data must participate even if it has no logical connection
- ⇒ But programmer is fully aware of non-local interactions and hence design to minimize
- ⇒ Can be efficiently implemented on a wide variety of architectures

I. Blocking + Non-Buffered

- Send does not return until matching receive is posted
- Message sent & send returns on completion of the communication operation
- Involves a handshake between sender & receiver
- No buffers used at either end



→ Process Idling


→ Best suited when send and receive are posted at roughly the same time

```
P0  
int a = 100, b = 200;  
send(&a, 1, 1);  
receive(&b, 1, 1);  
printf("f.d f.d", a, b);
```

```
P1  
int a = 200, b = 100;  
send(&b, 1, 0);  
receive(&a, 1, 0);  
printf("f.d f.d", a, b);
```

II . Blocking Buffered

- Assume pre-allocated buffers at sender and/or receiver
- `send(...)` copies the data into the designated buffer and returns
- Sender can continue - data change cannot influence semantics of communication
- Asynchronous communication capability of hardware
- At receiver, data is copied into a buffer
- When the destination process calls `receive(...)`, buffer is checked and if available, copies message to its buffer
- Almost NO Idling
- If such hardware is not available on both sides, some overhead can be avoided by using only one side buffers.
- On encountering a `send(...)`, the sender interrupts the receiver, both interact and message is deposited in a buffer at the receiver
- When receiver calls `receive(...)`, message is copied from the buffer into the target location



P₀

```
int a = 100, b = 200;
send(&a, 1, 1);
receive(&b, 1, 1);
printf(" %d %d", a, b);
```

P₁

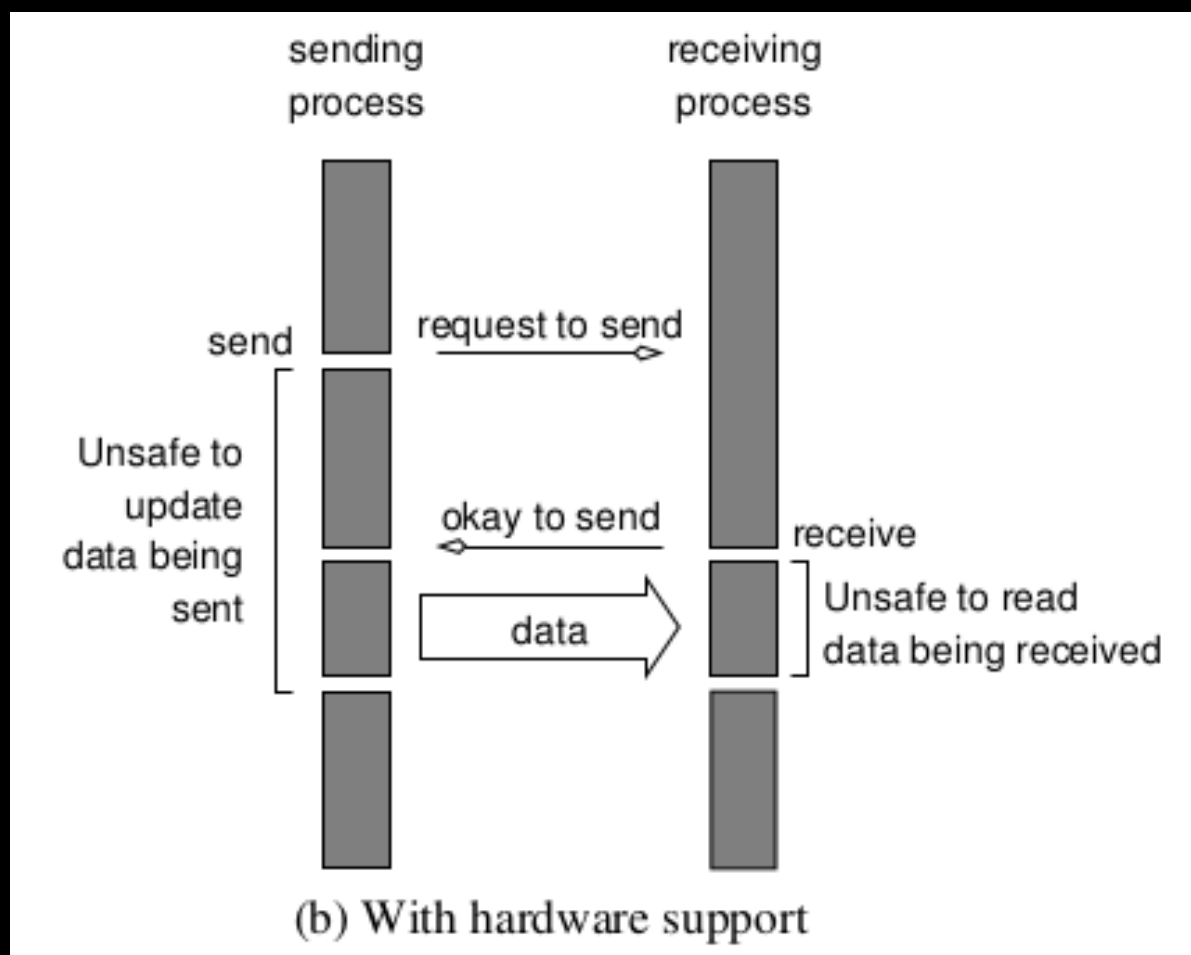

```
int a = 200, b = 100;
send(&b, 1, 0);
receive(&a, 1, 0);
printf(" %d %d", a, b);
```

P₀

```
int a = 100, b = 200;
receive(&a, 1, 1);
send(&b, 1, 1);
printf(" %d %d", a, b);
```

P₁

```
int a = 200, b = 100;
receive(&b, 1, 0);
send(&a, 1, 0);
printf(" %d %d", a, b);
```

IV. Non-Blocking + Buffered

- Sender initiates DMA and returns
- Data unsafe till DMA operation is complete
- Receiver initiates transfer from sender's buffer to the receiver's target location
- Reduces the time during which the data is unsafe

MP I : Message Passing Interface

- Standard library for Message passing
- Portable message passing in C or Fortran
- Defines syntax & semantics of the routines
- Almost all vendors provide their implementation
- > 125 functions
- Used in academia & industry
- Widely popular