

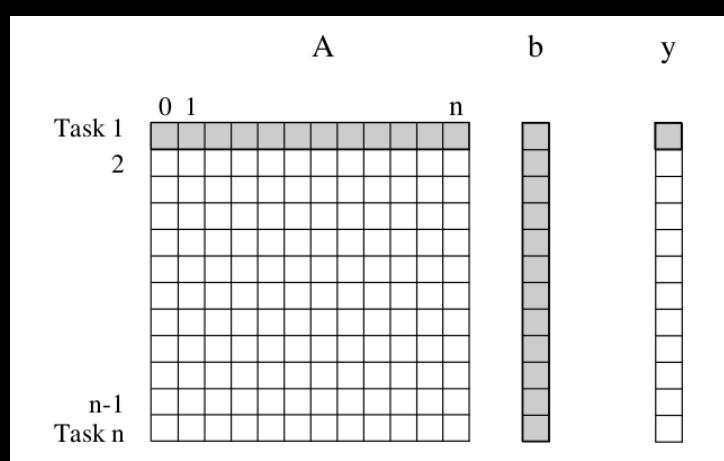
Steps involved to design (and implement) a parallel algorithm:

- Identify Concurrent Tasks
- Map Tasks to processes running in parallel
- Distribute input, output, intermediate data to processes
- Manage accesses to shared data by multiple processes
- Synchronize processes at various stages of parallel program execution

Example I: Dense Matrix-Vector multiplication

$$y[i] = \sum_{j=1}^n A[i][j] * b[j]$$

Task i =
Computation
of $y[i]$



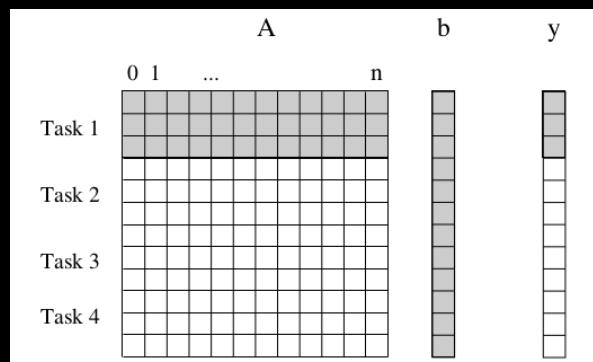
- All tasks are independent and hence can be executed simultaneously (any sequence)

Granularity : Number and size of tasks into which a problem is decomposed into.

→ Fine-grained - Large no. of small tasks

→ Coarse-grained - Small no. of large tasks

How much
more fine-
grained?
Any limit?



Maximum Degree of Concurrency : The maximum number of tasks that can be executed simultaneously in a parallel program at any given time.

Average Degree of Concurrency : Average number of tasks that can be run concurrently over the entire duration of execution of the program.

Granularity Vs Degree of Concurrency

→ Generally increases

→ Depends on task dependencies

Example II : Database Query Processing

Query : MODEL = "civic" AND YEAR = "2001" AND (COLOR = "Green" OR COLOR = "White")

ID#	Model	Year	Color	Dealer	Price
4523	Civic	2002	Blue	MN	\$18,000
3476	Corolla	1999	White	IL	\$15,000
7623	Camry	2001	Green	NY	\$21,000
9834	Prius	2001	Green	CA	\$18,000
6734	Civic	2001	White	OR	\$17,000
5342	Altima	2001	Green	FL	\$19,000
3845	Maxima	2001	Blue	NY	\$22,000
8354	Accord	2000	Green	VT	\$18,000
4395	Civic	2001	Red	CA	\$17,000
7352	Civic	2002	Red	WA	\$18,000

① MODEL = "Civic" \Rightarrow

ID#	Model
4523	Civic
6734	Civic
4395	Civic
7352	Civic

② YEAR = "2001" \Rightarrow

ID#	Year
7623	2001
6734	2001
5342	2001
3845	2001
4395	2001

③ COLOR = "White" \Rightarrow

ID#	Color
3476	White
6734	White

④ COLOR = "Green" \Rightarrow

ID#	Color
7623	Green
9834	Green
5342	Green
8354	Green

⑤ COLOR = "White"
OR
COLOR = "Green" \Rightarrow

ID#	Color
3476	White
7623	Green
9834	Green
6734	White
5342	Green
8354	Green

⑥ YEAR = "2001" AND
(COLOR = "White" OR
COLOR = "Green")

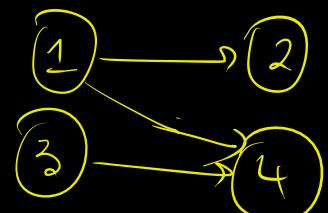
ID#	Color	Year
7623	Green	2001
6734	White	2001
5342	Green	2001

⑦ MODEL = "Civic" AND
YEAR = "2001" AND
(COLOR = "White" OR
COLOR = "Green")

ID#	Model	Year	Color
6734	Civic	2001	White

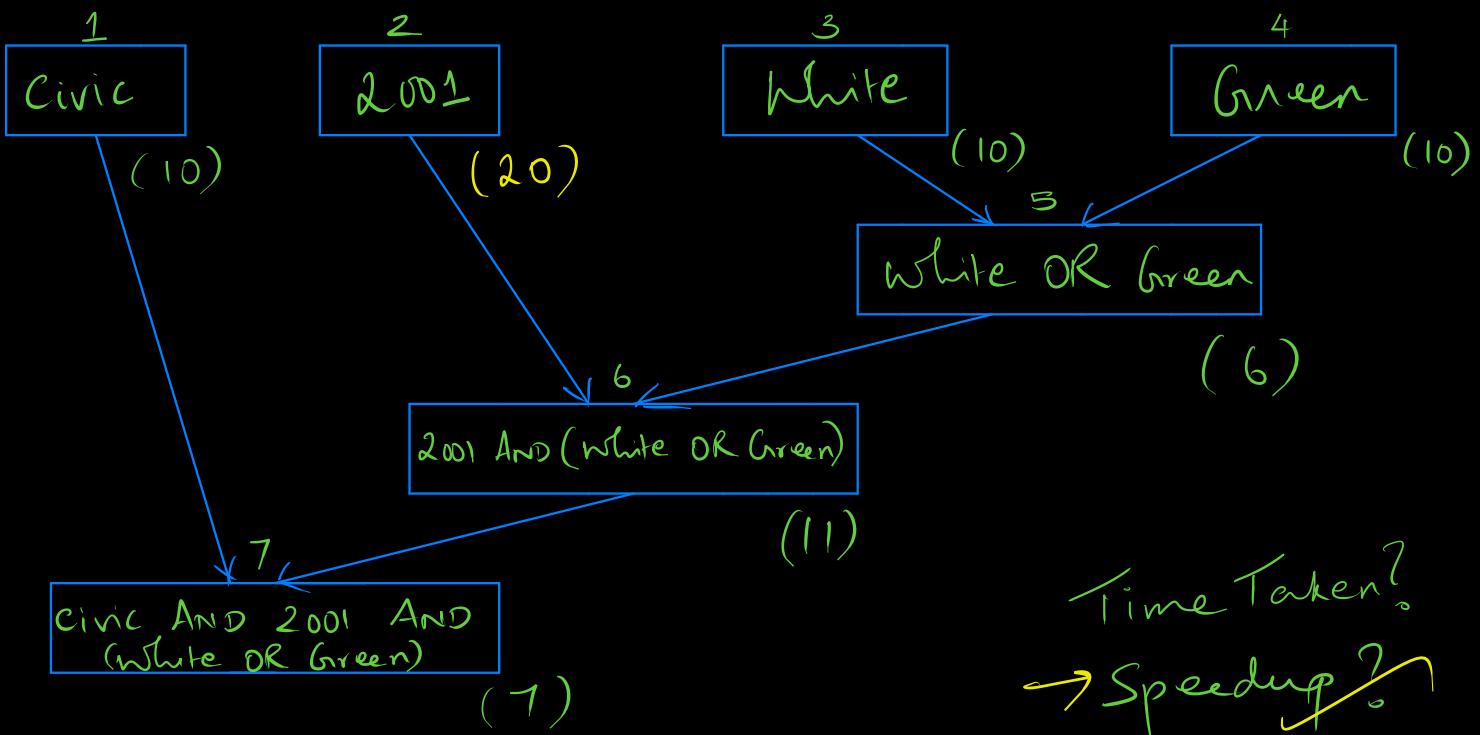
Task Dependency Graph: Directed Acyclic Graph in which nodes represent tasks and the directed edges indicate the dependencies amongst them.

→ Edge set can be empty



→ Can be disconnected

→ Task corresponding to a node can be executed when all tasks connected to this node by incoming edges have completed



Granularity : 7 tasks (coarse grained) No. of Processors?

Maximum Degree of Concurrency : 4

To find average degree of concurrency :

Start Node : Node with no incoming edges

Finish Node : Node with no outgoing edges

Critical Path : Longest ^{weighted} directed path between any pair of start and finish nodes.

Weight of a Node : Size or amount of work associated with the corresponding task.

Critical Path Length : Sum of weights of nodes along the critical path.

\Rightarrow Average Degree of Concurrency =
(Max Speedup)

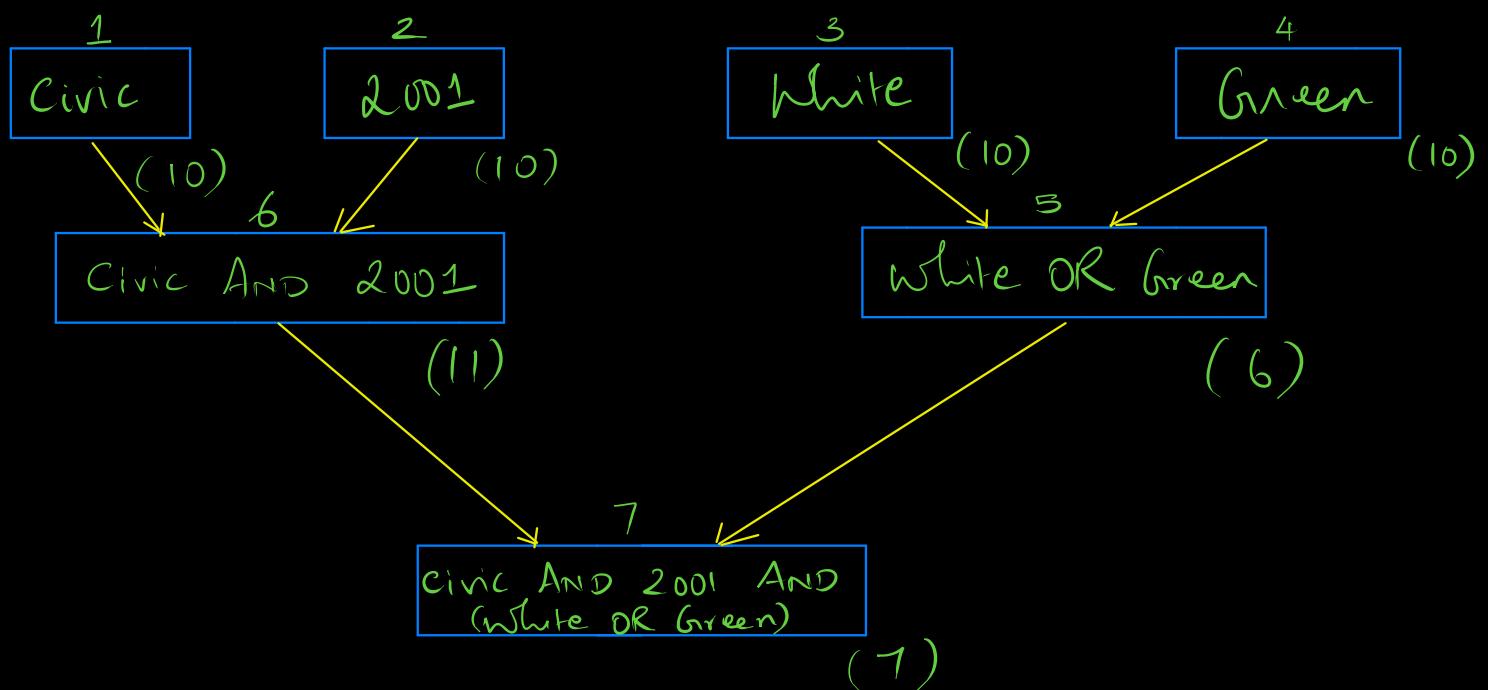
$$\frac{\text{Total amount of work}}{\text{critical path length}}$$

Time Taken?
→ Speedup?

\Rightarrow Shorter critical path favors a higher degree of concurrency.

$$\therefore \text{ADC} = \frac{64}{34} = 1.88$$

Better idea?



$$\text{Granularity} = 7$$

$$\text{Max Degree of Concurrency} = 4$$

$$\text{Average Degree of Concurrency} = \frac{64}{28} = \underline{\underline{2.28}}$$

Factors that limit speedup in parallelization:

→ Granularity

→ Degree of Concurrency (Task Dependency)

→ Task Interaction

Example III : Sparse Matrix Vector Multiplication

	A											b
Task 0	0	1	2	3	4	5	6	7	8	9	10	11
	●	●	●	●	●	●	●	●	●	●	●	●
	●	●	●	●	●	●	●	●	●	●	●	●
	●	●	●	●	●	●	●	●	●	●	●	●
	●	●	●	●	●	●	●	●	●	●	●	●
	●	●	●	●	●	●	●	●	●	●	●	●
	●	●	●	●	●	●	●	●	●	●	●	●
	●	●	●	●	●	●	●	●	●	●	●	●
	●	●	●	●	●	●	●	●	●	●	●	●
	●	●	●	●	●	●	●	●	●	●	●	●
	●	●	●	●	●	●	●	●	●	●	●	●
	●	●	●	●	●	●	●	●	●	●	●	●
	●	●	●	●	●	●	●	●	●	●	●	●
	●	●	●	●	●	●	●	●	●	●	●	●
	●	●	●	●	●	●	●	●	●	●	●	●

$$\begin{aligned}
 y[0] &= A[0][0] * b[0] + \\
 &\quad A[0][1] * b[1] + \\
 &\quad A[0][4] * b[4] + \\
 &\quad A[0][8] * b[8]
 \end{aligned}$$

$y = Ab$

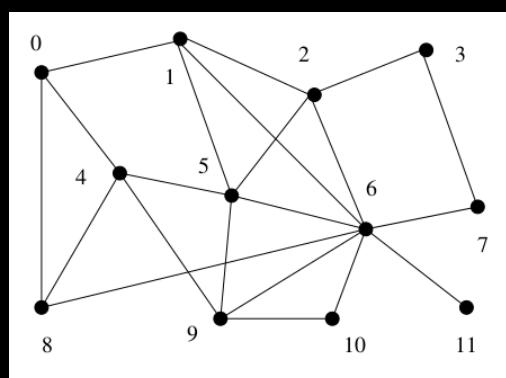
Partition Output Vector y and have each task compute an entry of $y \Rightarrow 12$ Tasks

Assume Task i is owner of $A[i][*]$ and $b[i]$
(input split across tasks)

But for Task i to compute $y[i]$, it needs other elements of b .

Task 8 needs $b[0]$ from Task 0,
 $b[4]$ from Task 4,
 $b[6]$ from Task 6.

Task Interaction
Graph :



(But each task is independent)

Granularity Vs Interaction :

Let time to Compute = 1 unit of time

∴ ∴ ∴ Communicate = 1 unit of time

⇒ Task 0 Computes in 1 unit of time
and communicates in 3 units of time

What if we assign Task 0 to Computing
elements 0, 4 & 8 of y ?

⇒ Task 0 already has with it
 $b[0]$, $b[4]$ and $b[8]$.

⇒ Needs 3 units of time to Compute.

How many units of time for communication?

4

Compute : Communication ratio 1 : 3

Which is favourable? _____

3 : 4

Process : Logical Computing agents that Perform Tasks.

→ Abstract entity that uses the code and data corresponding to a task to produce the output of that task within a finite amount of time after the task is activated by the parallel program.

Mapping : Mechanism by which Tasks are assigned to Processes for execution.

Good Mapping : [Task Dependency + Task Interaction graphs]

- Maximise Concurrency by mapping independent Tasks to different Processes T_1 T_2 P_0
- Minimise total completion time by ensuring Processes are available to execute Tasks on the critical path as soon as such Tasks are ready to be executed
- Minimise interaction among Processes by mapping Tasks with high degree of mutual interaction onto the same Process.

One Task
+
One Process

n Tasks
+
 n Processes

KEY : Finding a balance that optimizes Parallel Performance.

Decomposition \Rightarrow Degree of Concurrency

Mapping \Rightarrow Efficiency of the resulting Parallel Algorithm

